

Reproducing Sequence to Sequence Learning with Neural Networks

Sutskever, Vinyals & Le (Google), NeurIPS 2014 — arXiv:1409.3215

EE5180 course project · mid-term deliverable · TA: Prasenjit Kr Mudi (EE21D057)

1. What we set out to do

The mid-term brief is to **reproduce one or two rows of Table 1** and to demonstrate understanding of the method. We target rows 3 and 4 — a *single forward* LSTM against a *single reversed* LSTM — because that pair isolates the paper's central empirical claim: that reversing the word order of the source sentence, and not the target, markedly improves translation.


Table 1 row	Paper, cased BLEU on ntst14
Single forward LSTM, beam 12	26.17
Single reversed LSTM, beam 12	30.59

The scale gap — read this before any number in this report

The paper trains a **384M-parameter, 4×1000 LSTM on 12M sentence pairs across 8 GPUs for ~10 days**. We train a **2×512 LSTM on a 0.5M-pair subset on a single Apple M1 Pro**. Our absolute BLEU is therefore **not comparable** to the paper's, and we do not present it as such. What we claim to reproduce is the **direction, ordering and shape** of the paper's effects. Every table below repeats the relevant caveat.

2. The method, in the terms that matter for the experiment

Sequence-to-sequence with a reversed source (Sutskever et al., 2014, Fig. 1)




An **encoder** LSTM reads the source sentence one token at a time and ends in a state (h, c) . That state — and nothing else — is passed to a **separate decoder** LSTM, which is a language model conditioned on it.

Training maximises

$$\frac{1}{|S|} \sum_{(T, S) \in S} \log p(T | S) \quad \text{where} \quad p(T | S) = \prod_t p(y_t | y_{<t}, v)$$

and inference is a left-to-right beam search. Three design choices carry the paper's result: **separate encoder/decoder parameters**, **depth**, and **source reversal**.

Why reversal should help



Reversal leaves the *average* distance between corresponding words unchanged, but it collapses the **minimal time lag** — the gap between the first source word and the first target word it licenses — from the sentence length to one step. Backpropagation therefore has a short path to establish the source→target correspondence early in the sentence, and the paper argues this is what makes the optimisation problem easier (sec. 3.3).

The fixed vector **v** is also the paper's unresolved limitation. Every source sentence, however long, is squeezed into the same number of reals before a single target word is emitted. That bottleneck is what motivated attention, and it is the subject of our end-term half.

3. What we held fixed and what we changed

Reproduction is only meaningful if the deviations are declared, so here they are in full.

Held exactly as in the paper. Two separate encoder/decoder LSTMs with no shared parameters; the encoder's final (h, c) as the only conditioning signal; source reversed at train *and* test with the target never reversed; <E0S> terminating and scored; OOV → UNK; uniform initialisation U(−0.08, 0.08); plain SGD without momentum at lr 0.7, held constant then halved every half epoch; batches of 128; gradient-norm clipping at 5; length-bucketed minibatches; a naive full softmax; and a beam search in which a hypothesis leaves the beam as soon as it emits <E0S>, with candidates ranked by raw log-probability and no length normalisation.

One subtlety is worth stating because it is easy to get wrong. The paper defines clipping as “compute $s = \|g\|_2$ where g is the gradient divided by 128; if $s > 5$, set $g = 5g/s$ ”. We therefore normalise the loss by the **batch size**, not by the number of target tokens; normalising by tokens would silently change what the threshold of 5 means.

Declared deviations.

	Paper	Ours	Reason
Architecture	4 layers $\times$ 1000 cells, 1000-d embeddings, 384M params	2×512 , 512-d, $\sim$ 58M params	16 GB laptop
Sentence vector	8000 reals	2048 reals	follows from the above
Vocabulary	160k source / 80k target	32k / 32k	the naive softmax dominates cost
Training data	12M "selected" WMT'14 pairs	500,000 pairs from News-Commentary v9 + Europarl v7	compute budget
Training	7.5 epochs, 8 GPUs, $\sim$ 10 days	$\sim$ 8 epochs, 1 GPU, $\sim$ 5 h per run	compute budget
Ensembling	5 models	single models (seed ensembles on Tier 0 only)	wall-clock
Scoring	cased <code>multi-bleu.pl</code>	sacreBLEU <code>--tokenize none</code> (equivalent) and standard sacreBLEU	reproducibility

The 32k vocabulary leaves an OOV rate of **5.8%** on the English side of the test set and **5.9%** on the French side. The paper notes that its own BLEU was penalised on out-of-vocabulary words at 80k; ours is penalised considerably harder, and that is part of the scale gap rather than a separate defect.

4. Data and evaluation

Training. News-Commentary v9 (183,251 pairs) plus a seeded 1.1M-pair sample of Europarl v7 — both constituent corpora of the official WMT'14 En-Fr training set. After Moses tokenisation we drop empty pairs (3,207), pairs outside 1–50 tokens (170,275), pairs with a length ratio above 2.5 (1,823) and duplicates (20,066), leaving 1,087,880 pairs, from which we sample 500,000.

Test set. We evaluate on **newstest2014, full 3003-sentence version — the paper's own ntst14** — obtained through `sacrebleu -t wmt14/full -l en-fr`. Development is newstest2013. So our rows are scored on exactly the test set behind Table 1; only the model and the training data are smaller.

Two BLEU numbers, always both. The paper used cased `multi-bleu.pl` on *tokenised* text; modern practice is sacreBLEU on *detokenised* text. These are different measurements, so we report both throughout. The tokenised figure is the one comparable *in kind* to Table 1.

A note on data hygiene. News-Commentary v9 contains roughly 2,900 bare carriage returns *inside* lines — and a different number of them in the English and French files. Python's default universal-newline handling splits on those, which silently misaligns the parallel corpus. Our reader splits on `\n` only; two regression tests pin the behaviour. This is exactly the class of bug that produces plausible-looking but meaningless results, so we flag it rather than bury it.

5. Results — Table 1 rows 3 and 4

Scale gap. 0.5M training pairs against the paper's 12M; 2×512 against 4×1000; 32k/32k vocabulary against 160k/80k; one laptop GPU against eight datacentre GPUs for ten days. The absolute BLEU below is **not** comparable to the paper's column; the comparison that matters is forward against reversed within our own column.


Model	beam	BLEU (tokenized, cased)	BLEU (sacreBLEU, detok)	test perplexity	paper (Table 1)
Single forward LSTM	1	4.05	2.68	32.888	—
Single forward LSTM	2	4.24	2.86	32.888	—
Single forward LSTM	12	3.89	2.71	32.888	26.17
Single reversed LSTM	1	6.41	4.46	24.504	—
Single reversed LSTM	2	6.82	4.82	24.504	—
Single reversed LSTM	12	6.60	4.77	24.504	30.59

The central claim

- **BLEU at beam 12:** forward **3.89** → reversed **6.60**, a gain of **+2.71** BLEU. The paper reports 26.17 → 30.59, a gain of +4.42.
- **Test perplexity:** forward **32.89** → reversed **24.50**, a **25%** reduction. The paper reports 5.8 → 4.7, a 19% reduction.

The beam-size trend

Reversed model: B=1 **6.41**, B=2 **6.82**, B=12 **6.60**. BLEU **peaks at beam 2** and falls at beam 12, so the paper's monotone improvement with beam size does **not** reproduce at our scale. Beam 2 beating beam 1 is consistent with the paper; beam 12 beating beam 2 is not. The decomposition below shows why.



Why a wider beam hurts at our scale

The paper's BLEU rises with beam size. Ours peaks at beam 2. Decomposing BLEU shows why, and shows that length normalisation is *not* the fix.


ranking	beam	BLEU	brevity penalty	length ratio	1-gram prec	4-gram prec
raw log-probability (paper's setting)	1	6.41	0.979	0.979	36.07	1.44
raw log-probability (paper's setting)	2	6.82	0.966	0.967	36.39	1.69
raw log-probability (paper's setting)	12	6.60	0.943	0.945	34.75	1.75
length-normalised (alpha=1)	1	6.41	0.979	0.979	36.07	1.44
length-normalised (alpha=1)	2	6.82	1.000	1.024	35.07	1.63
length-normalised (alpha=1)	12	6.07	1.000	1.143	30.74	1.49

Reading it.

- With the paper's raw-log-probability ranking, widening the beam shortens the output: length ratio 0.979 -> 0.945 and the brevity penalty 0.979 -> 0.943. Every extra token lowers a sequence's total log-probability, so a larger beam prefers shorter hypotheses.
- Length normalisation removes the brevity penalty entirely (ratio 1.143, BP 1.000) but **does not fix BLEU**: it overshoots into over-long output and 1-gram precision falls 34.75 -> 30.74, so BLEU drops further (6.60 -> 6.07).

- So the cause is not calibration of length alone. A wider beam genuinely finds *higher-probability* hypotheses -- 4-gram precision rises with beam under the raw ranking -- but at this scale the model's likelihood and translation quality have diverged, so searching harder optimises the wrong thing. This is the documented 'beam search curse'; the paper's model, trained on 24x more data, was accurate enough not to suffer it.
- Roughly **16% of our output tokens are <unk>** (12,176 of 76,716), a direct consequence of the 32k vocabulary. The paper notes its own BLEU was penalised on out-of-vocabulary words at 80k; ours is penalised far harder, and that is a large part of the absolute gap.

Behaviour on long sentences (paper Fig. 3)




source length	sentences	forward BLEU	reversed BLEU	reversed / forward
1-10	405	5.67	9.53	1.68x
11-20	954	5.36	9.26	1.73x
21-30	832	4.03	7.19	1.79x
31-40	535	3.23	5.19	1.61x
41-inf	277	2.41	3.69	1.53x

Reversal helps at every sentence length — the reversed model is ahead in each bucket, so its overall advantage is not an artefact of short sentences.

But the paper's long-sentence robustness does not reproduce. The paper reports that "*the LSTM did well on long sentences*" (sec. 3.7, Fig. 3). Ours degrades from 9.53 BLEU on 1-10-token sources to 3.69 on 41+, and the forward model degrades the same way. That is consistent with a single fixed vector (2048 reals here, against the paper's 8000) being too small to carry a long sentence at this scale — the bottleneck hypothesis the end-term half of the project tests directly, by adding attention and rerunning this grid.

Training curves



Reversed: dev perplexity 80.73 (epoch 1) $\rightarrow$ 24.02 (epoch 8). *Forward*: dev perplexity 96.91 (epoch 1) $\rightarrow$ 30.25 (epoch 8).



6. Supporting evidence — Multi30k En $\rightarrow$ Fr

Before spending WMT compute we validated the entire pipeline on Multi30k (29k sentence pairs, same language direction), and it doubles as a seed-controlled replication of the reversal ablation.

This is not a Table 1 reproduction — different corpus, different domain (image captions), different test set, and a smaller 2×256 model. It is evidence that the *effect* survives at a scale where we can afford three seeds per arm.

Model	beam	BLEU (tokenized, cased)	BLEU (sacreBLEU, detok)	test perplexity	paper (Table 1)
Single forward LSTM	1	3.28	3.24	14.352	—
Single forward LSTM	2	3.52	3.48	14.352	—
Single forward LSTM	12	3.75	3.55	14.352	—
Ensemble of 3 forward LSTM	1	4.12	4.08	—	—
Ensemble of 3 forward LSTM	2	5.02	4.99	—	—
Ensemble of 3 forward LSTM	12	4.57	4.29	—	—
Single reversed LSTM	1	15.87	15.42	8.934	—
Single reversed LSTM	2	16.01	15.61	8.934	—
Single reversed LSTM	12	16.02	15.56	8.934	—
Ensemble of 3 reversed LSTM	1	18.56	18.05	—	—
Ensemble of 3 reversed LSTM	2	18.69	18.22	—	—
Ensemble of 3 reversed LSTM	12	18.58	18.08	—	—

Reversal moves tokenized BLEU **3.75** $\rightarrow$ **16.02** and test perplexity **14.35** $\rightarrow$ **8.93** at beam 12. The effect is proportionally *larger* than the paper's, which is what one expects: with only 29k pairs the forward model never learns a reliable source $\rightarrow$ target correspondence at all, so its output is fluent French that does not track the source. Reversal, by shortening the minimal time lag, is precisely what lets it start.



7. Engineering notes

Correctness gates before compute. `tests/test_smoke.py` asserts that beam search at `B=1` is exactly greedy decoding; that a wider beam never returns a lower-scoring hypothesis; that reversing the input by hand and setting the reversal flag cancel out; that padding does not change the encoder's sentence vector; that the model can overfit a two-sentence batch; that scoring the references against themselves yields BLEU

100; and that **MPS and CPU agree on loss and gradients to 1e-3** — PyTorch's Metal LSTM has a history of silently wrong numbers, and discovering that after an overnight run costs a day.

Where the compute went. Training runs on the GPU (MPS) at roughly 6,500 target words per second — coincidentally close to the 6,300 words/s the paper reports for its 8-GPU rig on a far larger model. Decoding, by contrast, runs on the **CPU**: beam search is latency-bound rather than throughput-bound, and Metal's per-kernel launch overhead made it **9× slower** than CPU (317 ms/sentence against 34 ms at B=12).

8. What this does and does not establish

Reproduced. Reversing the source improves both BLEU and perplexity by a clear margin under an otherwise identical setup, and the ordering of Table 1's rows 3 and 4 holds at our scale. The reversed model was ahead at *every* epoch, not merely at the end.

Not reproduced — and reported rather than hidden. The paper's monotone improvement with beam size does not survive at our scale: BLEU peaks at beam 2 and falls at beam 12. Section 5 decomposes this. Length normalisation removes the brevity penalty but makes BLEU *worse*, so the cause is not length calibration; a wider beam finds genuinely higher-probability hypotheses that are poorer translations, because at 1/24 of the paper's data the model's likelihood and translation quality have diverged.

Nor does the paper's robustness on long sentences: both of our models lose most of their BLEU past 40 source tokens, although reversal keeps its lead in every length bucket.

Not reproduced, and not attempted. Absolute BLEU anywhere near 26–31, the 5-model ensemble rows, and the 1000-best rescoring of Table 2. These need the paper's data and compute, and we say so rather than presenting a smaller number as if it were comparable.

9. Where the end-term goes

The single vector `v` is a fixed-capacity channel that every source sentence must pass through before decoding begins. The end-term half quantifies that bottleneck — BLEU and perplexity by source-length bucket, and a probe of encoder-state saturation — then implements **Bahdanau attention on this same codebase** and reruns the identical grid, so the fix is demonstrated rather than merely argued, before tracing the line from there to the Transformer.

Reproducing this report

```
make setup && make test
make data-multi30k && make smoke && make results-multi30k
make data-wmt && make wmt && make results-wmt
python scripts/make_report.py
```

Every number above is read from `results/*/all_results.json` and the training logs at build time, so the prose cannot drift from the runs.